

Lab 3 – Textual Data Analytics

Table of Contents

I. Feature Engineering	3
1. Text Normalization	3
2. Implement TF-IDF	3
3. Compare the results with the reference implementation of scikit-learn library.	4
4. Apply TF-IDF for information retrieval	4
II. Exercises	5

I. Feature Engineering

Complete the code with TODO tag in the Jupyter notebooks.

In this exercise we will understand the functioning of TF/IDF ranking. Implement the feature engineering and its application, based on the code framework provided below. Please read the file `feature_engineering.ipynb`

First, we use textual data from Twitter.

	id	created_at	text
0	849636868052275200	2017-04-05 14:56:29	b'And so the robots spared humanity ... https:...
1	848988730585096192	2017-04-03 20:01:01	b"@ForIn2020 @waltmossberg @mims @defcon_5 Exa...
2	848943072423497728	2017-04-03 16:59:35	b'@waltmossberg @mims @defcon_5 Et tu, Walt?'
3	848935705057280001	2017-04-03 16:30:19	b'Stormy weather in Shortville ...'
4	848416049573658624	2017-04-02 06:05:23	b"@DaveLeeBBC @verge Coal is dying due to nat ...

1. Text Normalization

Now we need to normalize text by stemming, tokenizing, and removing stopwords.

```
In [3]: ▶ def normalize(document):
# TODO: remove punctuation
text = "".join([ch for ch in document if ch not in string.punctuation])

# TODO: tokenize text
tokens = nltk.word_tokenize(text)

# TODO: Stemming
stemmer = PorterStemmer()
ret = " ".join([stemmer.stem(word.lower()) for word in tokens])
return ret

original_documents = [x.strip() for x in data['text']]
documents = [normalize(d).split() for d in original_documents]
documents[0]
```

As you can see that the normalization is still not perfect. Please feel free to improve upon (OPTIONAL), e.g. <https://marcobonzanini.com/2015/03/09/mining-twitter-data-with-python-part-2/>

2. Implement TF-IDF

Now you need to implement TF-IDF, including creating the vocabulary, computing term frequency, and normalizing by tf-idf weights.

```
In [4]: ▶ # Flatten all the documents
flat_list = [word for doc in documents for word in doc]

# TODO: remove stop words from the vocabulary
words = [word for word in flat_list if word not in stopwords.words('english')]

# TODO: we take the 500 most common words only
counts = Counter(words)
vocabulary = counts.most_common(500)
print([x for x in vocabulary if x[0] == 'tesla'])
vocabulary = [x[0] for x in vocabulary]
assert len(vocabulary) == 500

# vocabulary.sort()
vocabulary[:5]

[('tesla', 287)]
```

```
In [6]: ▶ def idf(vocabulary, documents):
        """TODO: compute IDF, storing values in a dictionary"""
        idf = {}
        num_documents = len(documents)
        for i, term in enumerate(vocabulary):
            idf[term] = math.log(num_documents / sum(term in document for document in documents), 2)
        return idf

idf = idf(vocabulary, documents)
[idf[key] for key in vocabulary[:5]]
```

3. Compare the results with the reference implementation of scikit-learn library.

Now we use the scikit-learn library. As you can see that, the way we do text normalization affects the result. Feel free to further improve upon (OPTIONAL),
e.g. <https://stackoverflow.com/questions/36182502/add-stemming-support-to-countvectorizer-sklearn>

```
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.metrics.pairwise import linear_kernel

tfidf = TfidfVectorizer(analyzer='word', ngram_range=(1,1), min_df = 1, stop_words = 'english', max_features=500)

features = tfidf.fit(original_documents)
corpus_tf_idf = tfidf.transform(original_documents)

sum_words = corpus_tf_idf.sum(axis=0)
words_freq = [(word, sum_words[0, idx]) for word, idx in tfidf.vocabulary_.items()]
print(sorted(words_freq, key = lambda x: x[1], reverse=True)[:5])
print('testla', corpus_tf_idf[1, features.vocabulary_['tesla']])

[('http', 163.54366542841234), ('https', 151.85039944652075), ('rt', 112.61998731390989), ('tesla', 95.96401470715628
1), ('xe2', 88.209444863464768)]
testla 0.349524310066
```

4. Apply TF-IDF for information retrieval

We can use the vector representation of documents to implement an information retrieval system.
We test with the query Q = "tesla nasa"

```
In [9]: ▶ def cosine_similarity(v1,v2):
        """TODO: compute cosine similarity"""
        sumxx, sumxy, sumyy = 0, 0, 0
        for i in range(len(v1)):
            x = v1[i]; y = v2[i]
            sumxx += x*x
            sumyy += y*y
            sumxy += x*y
        if sumxy == 0:
            result = 0
        else:
            result = sumxy/math.sqrt(sumxx*sumyy)
        return result

def search_vec(query, k, vocabulary, stemmer, document_vectors, original_documents):
    q = query.split()
    q = [stemmer.stem(w) for w in q]
    query_vector = vectorize(q, vocabulary, idf)

    # TODO: rank the documents by cosine similarity
    scores = [[cosine_similarity(query_vector, document_vectors[d]), d] for d in range(len(document_vectors))]
    scores.sort(key=lambda x: -x[0])

    print('Top-{0} documents'.format(k))
    for i in range(k):
        print(i, original_documents[scores[i][1]])

query = "tesla nasa"
stemmer = PorterStemmer()
search_vec(query, 5, vocabulary, stemmer, document_vectors, original_documents)
```

```
Top-5 documents
0 b'@ashwin7002 @NASA @faa @AFPAA We have not ruled that out.'
1 b'RT @NASA: Updated @SpaceX #Dragon #ISS rendezvous times: NASA TV coverage begins Sunday at 3:30amET: http://t.co/qxm0Dz4jPE. Grapple at ...'
2 b'Deeply appreciate @NASA's faith in @SpaceX. We will do whatever it takes to make NASA and the American people proud.'
3 b'Would also like to congratulate @Boeing, fellow winner of the @NASA commercial crew program'
4 b"@astrostephenson We're aiming for late 2015, but NASA needs to have overlapping capability to be safe. Would do the same"
```

We can also use the scikit-learn library to do the retrieval.

```
new_features = tfidf.transform([query])

cosine_similarities = linear_kernel(new_features, corpus_tf_idf).flatten()
related_docs_indices = cosine_similarities.argsort()[::-1]

topk = 5
print('Top-{0} documents'.format(topk))
for i in range(topk):
    print(i, original_documents[related_docs_indices[i]])
```

Top-5 documents
0 b'RT @NASA: Updated @SpaceX #Dragon #ISS rendezvous times: NASA TV coverage begins Sunday at 3:30amET: <http://t.co/qrm0Dz4jPE>. Grapple at ...'
1 b'Deeply appreciate @NASA's faith in @SpaceX. We will do whatever it takes to make NASA and the American people proud.'
2 b'@NASA Best of luck to the Cygnus launch'
3 b'RT @SpaceX: Success! Congrats @NASA on @MarsCuriosity!'
4 b'@ashwin7002 @NASA @faa @AFPAA We have not ruled that out.'

II. Exercises

By using the job market data, finish the following tasks to analyse the top important keywords for IT sector.

- ✓ Filter the jobs for IT sector only.
- ✓ Put the description of all jobs into a list.
- ✓ Use scikit-learn to get top 20 important keywords.
- ✓ Choose one favorite keyword and perform information retrieval with scikit-learn.